

# Class 08: Mini Project Breast Cancer

Nathan Joseph (PID: A17668656)

## Table of contents

Background . . . . .	1
Data Import . . . . .	1
Principal Component Analysis (PCA) . . . . .	3
4. Hierarchical Clustering . . . . .	11
Combining Methods . . . . .	12
Prediction . . . . .	14

## Background

In today's class we will be employing all R techniques for data analysis that we have learned thus far - including the machine learning methods of clustering and PCA - to analyze real breast cancer biopsy data.

## Data Import

Data from the Wisconsin Cancer CSV file:

```
fna.data <- "WisconsinCancer.csv"  
  
wisc.df <- read.csv(fna.data, row.names=1)  
  
head(wisc.df, 3)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.8	1001
842517	M	20.57	17.77	132.9	1326
84300903	M	19.69	21.25	130.0	1203

```

smoothness_mean compactness_mean concavity_mean concave.points_mean
842302      0.11840          0.27760        0.3001          0.14710
842517      0.08474          0.07864        0.0869          0.07017
84300903     0.10960          0.15990        0.1974          0.12790
symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
842302      0.2419           0.07871       1.0950        0.9053         8.589
842517      0.1812           0.05667       0.5435        0.7339         3.398
84300903     0.2069           0.05999       0.7456        0.7869         4.585
area_se smoothness_se compactness_se concavity_se concave.points_se
842302      153.40          0.006399       0.04904       0.05373        0.01587
842517      74.08           0.005225       0.01308       0.01860        0.01340
84300903     94.03           0.006150       0.04006       0.03832        0.02058
symmetry_se fractal_dimension_se radius_worst texture_worst
842302      0.03003          0.006193       25.38         17.33
842517      0.01389          0.003532       24.99         23.41
84300903     0.02250          0.004571       23.57         25.53
perimeter_worst area_worst smoothness_worst compactness_worst
842302      184.6            2019          0.1622        0.6656
842517      158.8            1956          0.1238        0.1866
84300903     152.5            1709          0.1444        0.4245
concavity_worst concave.points_worst symmetry_worst
842302      0.7119           0.2654         0.4601
842517      0.2416           0.1860         0.2750
84300903     0.4504           0.2430         0.3613
fractal_dimension_worst
842302      0.11890
842517      0.08902
84300903     0.08758

```

Removing the first column from the dataframe as this is the pathologist provided diagnosis which is the answer to if our cell samples are malignant or benign:

```
wisc.data <- wisc.df[,-1]
```

Creating diagnosis vector from the original dataset:

```
diagnosis <- wisc.df$diagnosis
```

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

```
[1] 569
```

There are 569 observations in this dataset.

Q2. How many of the observations have a malignant diagnosis?

```
count = 0
for (i in 1:nrow(wisc.data)){
  if (diagnosis[i]=="M"){
    count = count + 1
  }
}
count
```

```
[1] 212
```

There are 212 observations with a malignant diagnosis in this dataset.

Q3. How many variables/features in the data are suffixed with `_mean`?

```
sum(endsWith(colnames(wisc.data), "_mean"))
```

```
[1] 10
```

There are 10 features in the data that are suffixed with “`_mean`”.

## Principal Component Analysis (PCA)

The main function in base R called `prcomp()` we will use the optional argument `scale=TRUE` here as the data columns/features/dimensions are on very different scales in the original dataset.

Perform PCA on `wisc.data` by completing the following code:

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
attributes(wisc.pr)
```

```

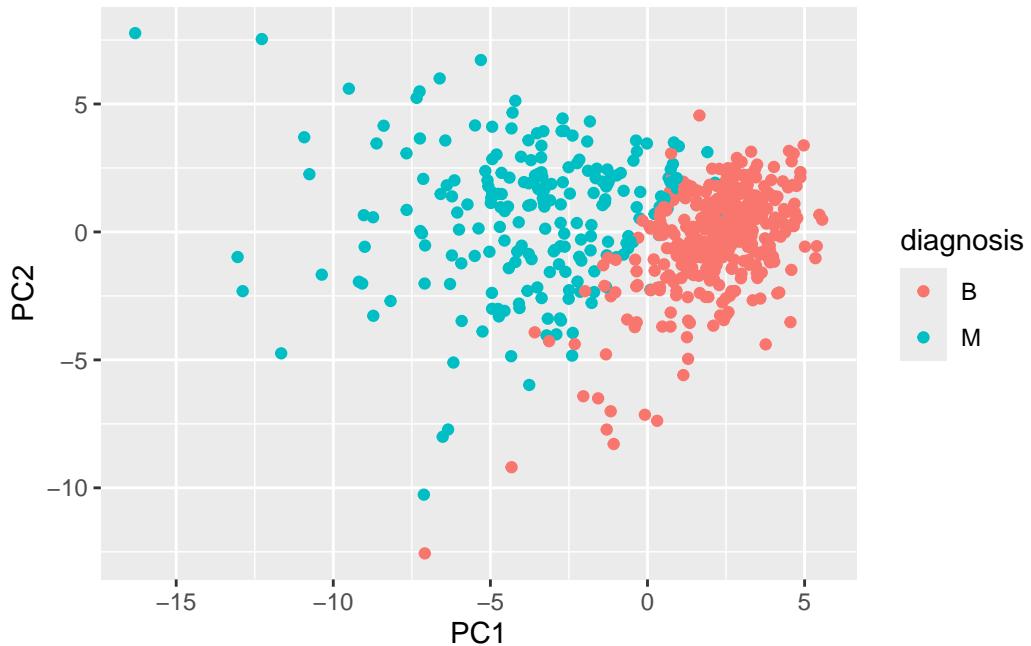
$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"

$class
[1] "prcomp"

library(ggplot2)

ggplot(wisc.pr$x) + aes(PC1, PC2, col = diagnosis) + geom_point()

```



Let's check the mean and standard deviation to see if scaling is necessary:

```
colMeans(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean
	1.412729e+01	1.928965e+01	9.196903e+01
	area_mean	smoothness_mean	compactness_mean
	6.548891e+02	9.636028e-02	1.043410e-01
	concavity_mean	concave.points_mean	symmetry_mean
	8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean		radius_se	texture_se
	6.279761e-02	4.051721e-01	1.216853e+00

perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst
1.323686e-01	2.542650e-01	2.721885e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
1.146062e-01	2.900756e-01	8.394582e-02

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

View the summary:

```
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

```
var_explained <- wisc.pr$sdev^2/sum(wisc.pr$sdev^2)
var_explained[1]
```

[1] 0.4427203

The first principal component captures 44.27% of the original variance.

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

```
var_explained <- wisc.pr$sdev^2 / sum(wisc.pr$sdev^2)
cumulative <- cumsum(var_explained)
which(cumulative >= 0.70)[1]
```

[1] 3

Three principal components are required to describe at least 70% of the original variance in the data.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

```
var_explained <- wisc.pr$sdev^2 / sum(wisc.pr$sdev^2)
cumulative <- cumsum(var_explained)
which(cumulative >= 0.90)[1]
```

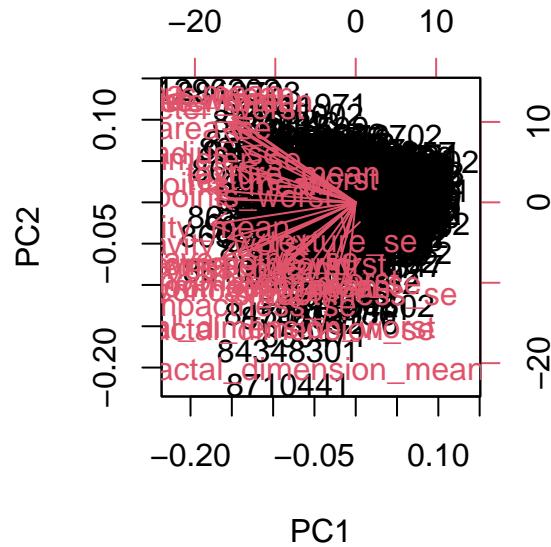
[1] 7

Seven principal components are required to describe at least 90% of the original variance in the data.

Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

A quick biplot of the principal components:

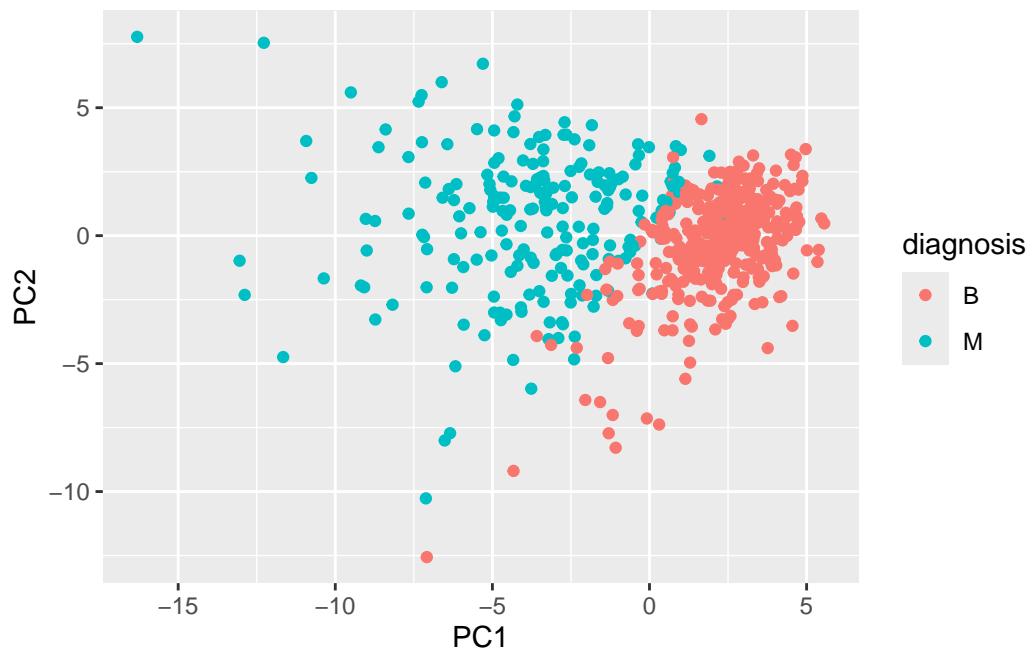
```
biplot(wisc.pr)
```



This plot shows a lot of features, dimensions, and a lot of numbers from the 2 principal components. Plotting PC1 vs PC2:

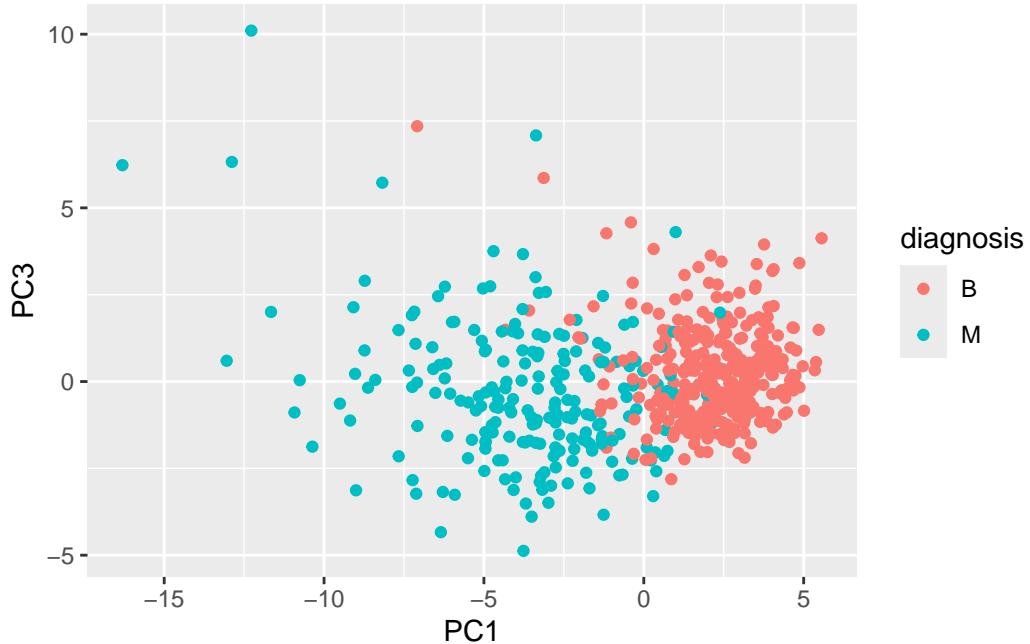
```
library(ggplot2)

ggplot(as.data.frame(wisc.pr$x)) +
  aes(PC1, PC2, col = diagnosis) +
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(as.data.frame(wisc.pr$x)) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()
```



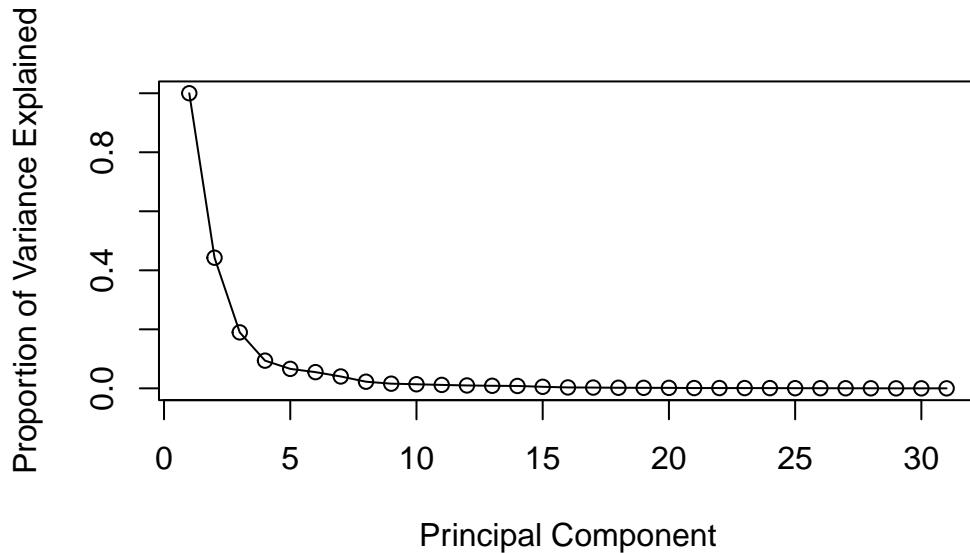
In these two plots, the principal component analysis shows that the principal components are identifying divisions in the data. In the above plot, the points are lower and more clustered towards the bottom.

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

Plotting the variance explained by each principal component:

```
pve <- pr.var / sum(pr.var)
plot(c(1,pve), xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[, 1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation[, 1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145
compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740
symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst	area_worst

```

-0.10446933      -0.23663968      -0.22487053
smoothness_worst compactness_worst concavity_worst
-0.12795256      -0.21009588      -0.22876753
concave.points_worst symmetry_worst fractal_dimension_worst
-0.25088597      -0.12290456      -0.13178394

```

The concave points means, which is a component of the loading vector, contributes the most to the `wisc.pr` data set in the first column. There are not any functions that are higher than the `concave.points_mean`.

## 4. Hierarchical Clustering

The goal of this section is to do hierarchical clustering of the original data to see if there is any obvious grouping into malignant and benign clusters.

```

data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled, method = "euclidian")
wisc.hclust <- hclust(data.dist, method = "complete")

```

```

wisc.hclust.clusters <- cutree(wisc.hclust, k = 2)
table(wisc.hclust.clusters)

```

```

wisc.hclust.clusters
 1   2
567  2

```

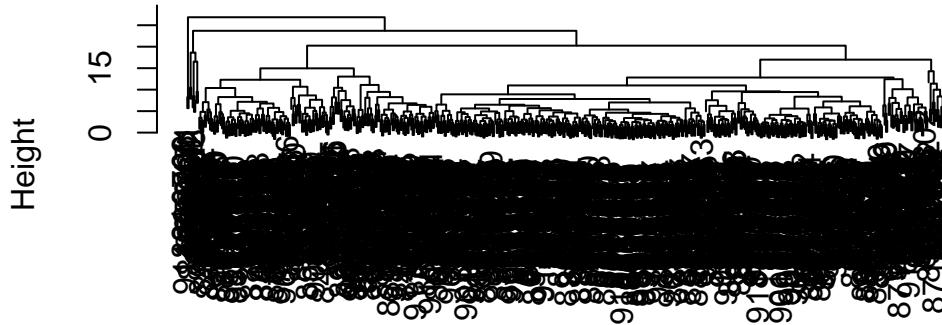
Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```

plot(wisc.hclust)
abline(wisc.hclust, col="red", lty=2)

```

## Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

The clustering model has 4 clusters at the height of 20.

Q12. Which method gives your favorite results for the same data.dist dataset?  
Explain your reasoning.

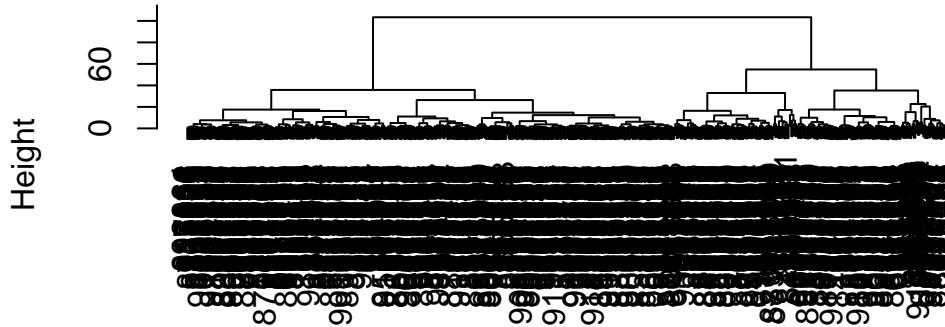
I prefer the cluster dendrogram as it is more detailed towards the original nodes of the dataset and also it best visualizes accuracy and variance in the dataset.

## Combining Methods

The idea here is that I can take my new variables (the PCs) that are better descriptors of the dataset than the original features (i.e. the 30 columns in wisc.data) and use these as a basis for clustering

```
pc.dist <- dist(wisc.pr$x[, 1:3])  
wisc.pr.hclust <- hclust (pc.dist, method = "ward.D2")  
plot(wisc.pr.hclust)
```

## Cluster Dendrogram



```
pc.dist  
hclust (*, "ward.D2")
```

```
grps <- cutree(wisc.pr.hclust, k = 2)  
table(grps)
```

```
grps  
 1   2  
203 366
```

I can now run `table()` with both my clustering `grps` and the expert diagnosis

```
table(grps, diagnosis)
```

grps	B	M
1	24	179
2	333	33

Q13. How well does the newly created `hclust` model with two clusters separate out the two “M” and “B” diagnoses?

The newly created `hclust` model with two cluster separates the two “M” and “B” diagnoses pretty well as demonstrated in more closed form dendrogram. The dendrogram is also much simpler.

Our cluster “1” has 179 “M” diagnosis Our cluster “2” has 333 “B” diagnosis  
179 TP 24 FP  
333 TN 24 FN Sensitivity: TP/(TP + FN)

```
179/(179+33)
```

```
[1] 0.8443396
```

Specificity: TN/(TN+FP)

```
333/(333+24)
```

```
[1] 0.9327731
```

## Prediction

We can use our PCA model for prediction of new un-seen cases.

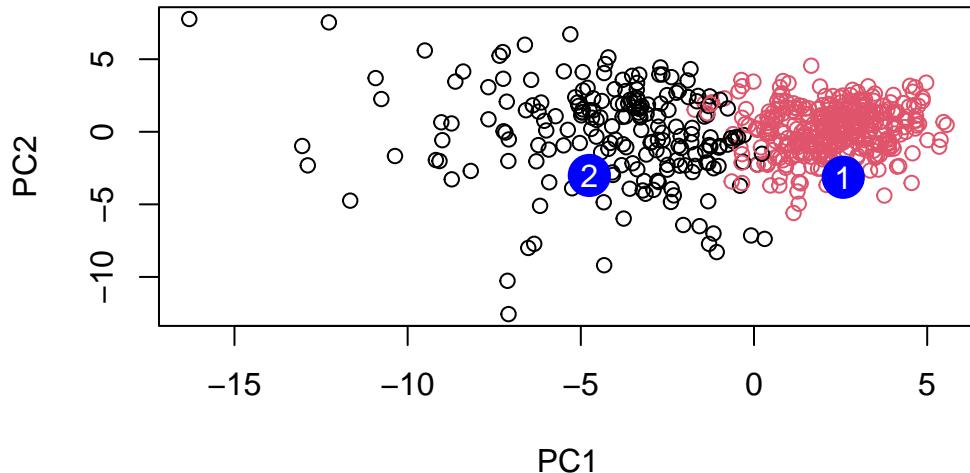
```
url <- "https://tinyurl.com/new-samples-CSV"  
new <- read.csv(url)  
npc <- predict(wisc.pr, newdata=new)  
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

```

plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q16. Which of these new patients should we prioritize for follow up based on your results?

Since there is a tighter cluster around patient 1, this is the patient we should prioritize for a follow up based on the results as there is a large variation or spread around patient 2.